SLURM and Alpine HPC Command Cheat Sheet

General SLURM Commands

The following are commands that will work on any HPC that uses SLURM as its job scheduler.

Command	Usage	Details and Options
`sbatch`	`sbatch nameOfMyScript.sh`	`sbatch -h` to see all options you can add. Notice, if you didn't want to add the `#SBATCH` directives to your script, you can pass them through the command line at run time
`squeue`	`squeue -u $USER`	Shows all the jobs you have submitted, are running, or have completed/finished in the 24 hour day and what their status is in the queue
`sstat`	`sstat -j 2793287 -o JobID,AveRSS,MaxRSS,MinCPU`	Only works on actively running jobs! Note you cannot get the Elapsed field using `sstat`. Running `sstat --helpformat` will tell you all of the possible outputs you can get statistics on for the job; Note, I have only output `JobID,AveRSS,MaxRSS,MinCPU` because I find those the most useful

`sacct`	`sacct`
	or
	`sacct -j 2793287.batch -o JobID,Elapsed,AveRSS,MaxRSS,MinCPU`
	or
	`sacct -j myJobID -B`
		If you want a summary of all jobs and their statuses in the 24 hour day, you can run `sacct` by itself; if you provide `sacct` with a job id after a job has completed/finished, it will give you various statistics about the job. See `sacct -h` to see different ways of filtering jobs. The `-B` option prints out the exact sbatch shell script you submitted for that job id.

`sinfo`	`sinfo`	Will give you a list of all the partitions on the system, what their default time limits are and the node identifiers that belong to that partition

Alpine Specific Commands

The following commands are specific to Alpine HPC and may not work on other clusters.

Command	Usage	Details and Options

`acompile`	`acompile`	This gives you interactive access to a compute node. A common option will be to change the time to have it longer. The default is 1 hour, but you can request up to 12 hours. Ex: `acompile --time=12:00:00`
`curc-quota`	`curc-quota`	If using from compute node, first run `module load curc-quota` This will tell you how much space you have and have used in your home, projects, scratch, and any petalibrary allocations
`levelfs`	`levelfs $USER`	If using from login/head node first run `module load slurmtools` This tells you how your job is prioritized and what your fairshare is.
`seff`	`seff myJobID`	This will tell you CPU and memory efficiency of a job that has already finished running
`jobstats`	`jobstats $USER pickDays`	If using from login/head node first run `module load slurmtools` Shows all jobs run over the last x days. For example: `jobstats $USER 5` would show me all the jobs I submitted over the last 5 days.